

OpenWrt port on the Nubia RedMagic 7 (NX679J)

Analysis report on the *openwrt-nx679j* repository: hygiene, reproducibility, phases, AI agents and models, human/AI attribution and cost estimate

Repositories analysed: github.com/CAPP-TESTS/openwrt-nx679j and github.com/Saddytech/openwrt-nx679j (commit f9a92fc)

Report date: 4 October 2026

Method: cloning and local analysis of the entire tree (2,520 files), plus verification of the public model prices

Summary

The repository documents a port of OpenWrt 25.12.5 to the Nubia RedMagic 7 (NX679J, Snapdragon 8 Gen 1 / SM8450). The port is not “native” OpenWrt in the strict sense: it is an OpenWrt userland running in a chroot on slot B, on top of the Android vendor kernel 5.10.66, which is left untouched. The X65 modem is driven directly over QMI/QRTR, without ModemManager, and the interface is a UI written in C on DRM/KMS.

The work was carried out almost entirely by AI agents, with a human role of direction, supervision and physical intervention on the device. The main findings of the analysis are the following.

- **Good hygiene, weak reproducibility.** No sensitive personal data is exposed and secrets are properly redacted. However, the build scripts for v91-v180 are missing, all the scripts present use absolute paths from the author’s machine, and there are no procedures for the OpenWrt rootfs or for flashing.
- **Four phases.** A pre-Hermes phase (2 July – 7 August) of unknown provenance, a model-testing phase (4–15 September), the actual Hermes phase (16–29 September, about 180 build iterations) and publication (30 September).
- **Not just Hermes + DeepSeek.** DeepSeek does not appear in any file. The documented models are GPT-6 Astra (design and diagnostic consultant), gpt-5.6-luna (testing), claude-haiku-4-5 (failed attempt) and OpenRouter (failed). Image analysis requires a model with visual input, different from DeepSeek V4 Pro. The P0 phase bears the stylistic fingerprint of a different, unidentifiable agent.
- **Attribution.** 97–99.5% of the volume produced and 93–98% of the execution are agentic. On the direction of the decisive turning points, however, the estimated human share is 25–40%.
- **Estimated costs.** About \$67 in the minimum scenario, \$176 in the medium and \$523 in the maximum. The most realistic estimate is about \$190–200, with a defensible range from about \$35 to about \$800. If the main model were GPT-6 Astra instead of DeepSeek, the main loop alone would cost about \$2,500.

Methodological note. The git history has been squashed into a single commit and the Hermes session logs (state.db, request_dump_*.json) are not published. All attributions and volumes are therefore reconstructed from the content of the files. The figures should be read as reasoned estimates, not measurements: each estimate comes with the assumptions it depends on.

1. The project

Stated goal. Boot OpenWrt on a high-end Android smartphone, obtain a persistent control channel over the single USB-C port, use the cellular modem as uplink and always keep a recovery path: slot A with stock Android and Magisk, or Qualcomm EDL mode. The README presents it as a “complete port” of OpenWrt 25.12.5.

Actual architecture. In the documents, “native” means that OpenWrt runs directly on the hardware with no Android running, but the kernel remains the manufacturer’s. The project does not port its own or a mainline kernel. The attempt with an upstream 6.6.y kernel in phase P0 failed: it hung after 15 seconds for lack of support for the Gunchah hypervisor. The final result is therefore an OpenWrt userland grafted onto the existing Android platform.

Component	Adopted solution	Notes
Kernel	Vendor 5.10.66 (android12), untouched	Vendor modules loaded from the ramdisk, no kernel patches
Boot	Custom ramdisk in the slot B boot image (header v4)	Own PID 1 and init (init v8/v9), USB NCM gadget on 10.0.0.1
Userland	OpenWrt 25.12.5 (r33051-f5dae5ece4), armv8, musl	Run in a chroot, with LuCI, SSH and persistence inside the ramdisk
Modem	Qualcomm X65 via QMI over QRTR (AF_QIPCRTR)	ModemManager retired from v130: own chain, LuCI plugin and link supervisor
Wi-Fi	QCA6490 / WCN6855 (cnss2)	NX679J-TEST access point at 5 GHz, NAT to the cellular WAN
Display and input	Native C UI on DRM/KMS, continuous atomic commits	The panel collapses without a commit for more than ~58 ms; navigation with the side keys
Recovery	Slot A Android + Magisk, fastboot, EDL	Slot switching is autonomous via a reboot binary and fastboot

Repository identity. CAPP-TESTS/openwrt-nx679j and Saddytech/openwrt-nx679j contain the exact same commit (f9a92fc, 30 September 2026 15:46 +0200, message “first commit”, author Saddytech). Release v1.0 was published on 1 October. The CAPP-TESTS repository should therefore be considered a copy, or an unmodified fork, of Saddytech’s.

Stated status over time. The technical report of 20 September explicitly stated “Internet NOT yet reached”: the first cellular ping caused a reboot. Later documents record the data connection with the standard proto (22 September), the elimination of the hourly bearer limit after 72 minutes of verification (24 September), the modem without ModemManager (v130) and the final display standby (v180, 25 September). These claims are documented but cannot be independently verified from the repository, because the images and the final builds are missing (see section 2).

Repository structure.

Folder	Content
01-documenti	Handoffs, session prompts, reports, snapshots of STATO-ATTUALE and DOCUMENTAZIONE
02-sorgenti	Init v8/v9, relay/worker, build-v54...v90.py, C UI, patches, live configurations
04-persistenza	Sanitised etc.tar, luci.tar, tools.tar, as injected into the ramdisk
05-bootchain-re	Boot-chain reverse engineering, native baseline, bootloader-re (phase P0), research on a CVE
06-ricerche	port-work (phase P0), research on DSI/display, idlepc
07-evidenze-runtime	Properties, dmesg, GPIO and regulators collected from Android on 11 July

Folder	Content
08-skill	18 Hermes skills (method, diagnostics, user rules, research fan-out)
09-albero-originale	Rest of the working tree: EDL, headless builds, hypervisor-bypass, subagent logs




Figure 1 — Repository composition by file type. Only 21% of the files are material written by the project (code and prose).

2. Repository hygiene and reproducibility

2.1 Hygiene

On this front the repository is well kept. The author's automatic verifier reports 23 out of 23 checks passed: secrets searched even inside binaries and tar archives, personal data, vendor artefacts recognised by extension and magic bytes, consistency of the stated figures. Gitleaks reports 2 findings, documented as false positives. I independently verified the most critical points.

Check	Result	Detail
Radio identifiers (IMEI, IMSI 22288..., ICCID 89..., MSISDN +39)	No occurrences	Pattern search across the entire tree
Redactions	≈ 9,000 markers	Serial number masked (3dbd****), Wi-Fi key [REDACTED]
OpenWrt credentials (etc.tar)	Clean	Root password locked (!), authorized_keys sanitised
IP addresses	Private only	10.0.0.1 (USB gadget), 10.0.2.x (QEMU), one carrier CGNAT address
Licences	Present	MIT for code, CC BY 4.0 for documents, LICENSES/ and THIRD_PARTY_LICENSES.md
Proprietary artefacts	Excluded by design	Firmware, .ko modules, boot images and Firehose loaders are not included

There are, however, also limitations that concern transparency rather than security.

- **Squashed history.** A single commit with 2,520 files: it is impossible to reconstruct who changed what and when.
- **Invasive deduplication.** DEDUPLICA.md declares the removal of 1,450 files (about 59 MB) from 239 folders. For some files it names as canonical copy a “revision archive” that is not published, for example several versions of STATO-ATTUALE.md and the ESITO and DECISIONE files.
- **The verification covers hygiene, not functionality.** None of the 23 checks demonstrates that the port can be rebuilt or booted.

2.2 Reproducibility

The README states that all excluded artefacts “can be rebuilt” and that ESTRAZIONE-BLOB.md contains “complete and verified” procedures. Verification gives a more nuanced result: the procedures to extract the vendor blobs are there, but the rest of the chain is incomplete.

Requirement for rebuilding	Status in the repo	Assessment
Stock firmware, partitions, modules, touch, DTB/DTBO	Procedures present (sections 1-6 of ESTRAZIONE-BLOB.md, 234 lines)	Sufficient for an expert
Firehose loader, GPT, EFS, GPL sources	Procedures present (sections 7-10)	Sufficient; EDL, however, “worked” only in the past
OpenWrt 25.12.5 rootfs	No procedure in ESTRAZIONE-BLOB (0 occurrences); only the source URL in DOCUMENTAZIONE	Gap
Image build scripts	38 scripts: v10 and v54-v90. v91-v180 are missing, and the documented final version is v180	Serious gap
Script portability	All 38 contain absolute paths /home/user/nx679j-stock/...	Must be adapted by hand

Requirement for rebuilding	Status in the repo	Assessment
Flashing procedure	0 occurrences of flash commands in README and ESTRAZIONE-BLOB	Gap
Persistence (etc/luci/tools)	Present, in sanitised form	Usable
UI, init, relay, QMI chain sources	Present (nx679j-ui.c, init v9, qmi-qtr.c, ...)	Usable
Ready-made images or rootfs	Absent by design	Consistent with the licence, but not bootable as is

Verdict. The port is **not reproducible “turnkey”**. An expert reverse engineer with the device could probably rebuild a version around v90 starting from build-v90.py: they would need to obtain the blobs with ESTRAZIONE-BLOB, download the OpenWrt rootfs, fix the paths and work out the flashing procedure from the logs. The final state described in the documentation (v180, with full UI and modem without ModemManager) cannot be rebuilt from the repository alone, because the builders for the last 90 iterations are missing.

3. Phases of activity

Dating cannot rely on metadata: git contains a single commit and all files have the same date. For about half of the files I derived the date from the path or the content (patterns such as 20260917 or 2026-07-10). The others I assigned by folder, following the chronology that DOCUMENTAZIONE.md itself reconstructs. Precision is therefore good at phase level, not at the level of individual days.




Figure 2 — Phase timeline.

Phase	Period	Main activities	Agent
P0	2 July – 7 August (plus the handoff of 16/09 at 10:30)	Download of the V311 OTA, Magisk patch, slot B backup, OpenWrt prototypes phases 1–4, stock modules, EDL readback of 17/07, RE of the bootloader and Gunyah, failed attempts with upstream kernel 6.6.y	Unknown; not Hermes
P0b	4–15 September	04/09: 7 model and provider test sessions (gpt-5.6-luna, claude-haiku-4-5 failed). 14/09: .hermes folder in hypervisor-bypass, with no session recorded	Hermes (configuration)
P1	16–29 September	Boot chain, init/PID 1, USB gadget, OpenWrt chroot, QMI/QRTR, Wi-Fi, LuCI, DRM/KMS UI, standby, boot-time optimisation: v10 → v180	Hermes + subagents + GPT-6 Astra + vision model
P2	30 September – 1 October	Writing of DOCUMENTAZIONE.md (256 KB), deduplication, checks, licences, single commit and release v1.0	Hermes (multiagent-document-assembly skill)




Figure 3 — Image versions per day, from the timeline table in DOCUMENTAZIONE.md. The 18/09 value is inferred from the jump v16 → v35.

The pace is intense: about 170 versions in nine days, with sessions documented at night as well (00:16–01:23, until 07:15) and a peak of 40 versions on 25 September, devoted mainly to the UI

and display standby.




Figure 4 — Lines of code and prose written by the project per phase, excluding third-party sources and tool output.

Phase	Files written	Lines	Stylistic signature
P0 pre-Hermes	67	~10,500	English (22 documents out of 24), 97 emoji 🐞🐛, “Executive Summary”, “Root Cause”
P1 Hermes	460	~81,300	Italian, 178 FATTO/IPOTESI/DA CONFERMARE markers (fact/hypothesis/to be confirmed), agent first person
P2 publication	4	~7,700	DOCUMENTAZIONE, DEDUPICA, README, ESTRAZIONE-BLOB

4. AI agents and models across the phases

The starting hypothesis was an orchestration made up only of Hermes and DeepSeek Pro. The analysis confirms Hermes but refutes the exclusivity: at least three other models appear in the work, plus an unidentified agent in phase P0.

Agent / model	Phase	Role	Evidence in the repo	Status
Hermes Agent (Nous Research, MIT)	P0b-P2	Orchestrator, kernel-re profile	deleg_* delegations, 7 subagent transcripts, skills signed “Hermes Agent”	Confirmed
Hermes subagents	P1	9-10-agent research fan-outs, batch analyses	task-b1-*/b2-* transcripts: 321 tool calls in total, no human intervention	Confirmed
DeepSeek (V4 Pro)	?	Hypothesis: main model	0 occurrences in the whole tree	Not found
GPT-6 Astra (OpenAI) via “oneprovider”	P1	Consultant: display diagnosis (RCG timeout), “DESIGN approvato da Astra” for the persistent DRM client	Command hermes -m gpt-6-astra --provider oneprovider --reasoning max; one modem consultation timed out	Confirmed
OpenRouter	P1	Alternative channel for Astra	HTTP 402 credit error, then explicitly forbidden	Failed
gpt-5.6-luna (OpenAI)	P0b (04/09)	Configuration test	First recorded Hermes session: 8 messages, no work on the device	Confirmed
claude-haiku-4-5	P0b (04/09)	Configuration test	404 model_not_found error	Failed
Vision model for vision_analyze	P1	Reading frames from the webcam pointed at the phone	DeepSeek V4 Pro does not accept image input, so another model is required	Unknown
P0 phase agent	P0	Bootloader RE, prototypes, reports in English	Different style, FINDINGS → SOLUTION_FOUND → REALITY_CHECK cycle	Unknown

4.1 The fingerprint of phase P0

The documentation itself admits that the work of 2 July – 7 August “is not attributable to this Hermes profile nor to any recorded tool”. That it was nonetheless AI-assisted emerges from several clues.

- In bootloader-re/ there is a sequence of files from 6 August with names typical of an LLM’s self-correction: FINDINGS, SOLUTION_FOUND, REALITY_CHECK, REAL_SOLUTION_DEEP_ANALYSIS, ACTUAL_RESULTS. FINDINGS attributes kvm-arm.mode=protected to the ABL; SOLUTION_FOUND contradicts it (the parameter would be in CONFIG_CMDLINE) without withdrawing the previous report.
- REALITY_CHECK addresses the user directly (“all your working OpenWrt builds”): this is the voice of a chat assistant.
- hyp_attack.c uses addresses commented as “SM8350 Gunyah addresses” on an SM8450 device: a typical error of code generated without verification.
- FINAL_REPORT states “Framework: reverse-skill + radare2”. Reverse-skill, however, is a client-agnostic skill pack, compatible with Claude Code, Kiro, Cursor, Cline and others: it does not identify the tool used.




Figure 5 — Stylistic markers compared between P0 and P1.

4.2 A false positive to discard

The string “Assisted-by: Claude:claude-sonnet-4-6” appears in tools/qdlrs/AGENTS.md. That file, however, is the policy of Qualcomm’s upstream qdl-rs project, included as a dependency: the string is the format example for the “Assisted-by” trailer under the Linux kernel convention, not evidence that Claude was used in the port.

5. Distribution of human and AI involvement

The ratio between human and AI work changes a lot depending on what is measured. I therefore used three distinct metrics.

5.1 The documented human footprint

The file `08-skill/nx679j-openwrt/references/regole-utente.md` collects, almost verbatim, the requests the user made in the 16–22 September sessions, and notes that they were “binding” because the user “repeatedly got irritated when they were not applied”. Together with the other quotes scattered across the repo, there are about thirty sentences actually spoken by the user. Some representative examples (translated from Italian):

“but before writing any more code, as I’ve told you 2 million times, look for the info online, unleashing even 20 agents at once”

“we were supposed to use standard OpenWrt components, not create our own protocol that serves no purpose”

“I’ve connected a webcam that looks at the smartphone, why don’t you find a way online to use it so that you yourself have visual input without my support?”

“and now you want to tell me you can’t get in via ssh either from the gadget or from wifi!? come on! I can see the gadget connected and I can see the phone’s wifi network!”

Type of human intervention	Documented examples
Direction and method	Search before writing code (origin of the block-research-fanout skill); use standard OpenWrt components (custom proto removed); persistence; Android as an “oracle”; no false success claims
Knowledge of the environment	Wi-Fi IP 192.168.77.1; fastboot → slot A → Magisk procedure; challenge of the “unreachability”; correction of a wrong visual reading (“cyan line”)
Physical actions	Connecting and aiming the webcam; reboots and power cycles (“I do the reboots”, at most one user action per test); side-key verification; reporting glitches
Authorisations	23/09: only one real DRM commit allowed, with stated limits (no flashing, no modem changes)
Writing code or documents	No trace; even the skills co-signed “user, Hermes Agent” are written by the agent

Agentic side. The 7 subagent transcripts contain 321 tool calls and zero human interventions, at a pace of about 7 tool calls per minute. The agent also had an autonomous visual verification loop (webcam, OBS and `vision_analyze`), which reduced the need for human input. That loop, however, came from an explicit request by the user.

5.2 The ranges




Figure 6 — Estimated agentic share per metric; the complement is the human share. The dashed line marks the stated “over 90%” threshold.

Metric	AI share	Human share	Rationale
Volume produced (lines of code and documents)	97-99.5%	0.5-3%	~99,500 lines written by the agent; direct human input is almost only in the quoted sentences. Margin for possible untracked manual edits in P0
Execution (actions on the system and the device)	93-98%	2-7%	Thousands of agentic commands and tool calls versus tens or a few hundred physical actions; many reboot and slot-switch cycles were autonomous
Direction (decisive turning points)	60-75%	25-40%	Several methodological turns, the observation infrastructure and various factual corrections come from the user. Qualitative estimate

Conclusion. The “over 90% agentic” figure is correct, even conservative, if it refers to what was written and executed. It does not describe well who steered the project: there the human role is a minority but far from marginal. Without the user the agent had got stuck in passive waits, had rewritten components that already existed and had declared a connected device unreachable.

Limitations. The quotes preserved are those the agent chose to transcribe, presumably the most pointed ones, so the actual human turns are underestimated. With state.db, the direction metric would become a count.

6. Cost estimate: three scenarios

There are no consumption logs in the repository, so token volumes are reconstructed from the artefacts. Prices, on the other hand, are the public ones in force in September–October 2026. For the main loop, DeepSeek V4 Pro is assumed, per the initial hypothesis: the main model is not declared in the repo, and this is the real unknown of the estimate.

6.1 Volume assumptions (scenario with all activities)

Component	Assumption	Basis in the repo
Hermes main loop	~10,000 turns; average context ~90k tokens; 90% cache hit; ~1,500 output tokens per turn	~180 builds, work from 16 to 30/09 including nights; status documents of ~30k tokens re-read often
Subagents	~175 subagents × 25 turns; context ~35k; 75% cache hit; ~1,000 output tokens per turn	Pace measured in the 7 transcripts; 9–10-agent fan-outs; 12 RICERCA/ESITO files; 4 delegation IDs
Orchestration overhead	+8% on the main loop	Context compression, memory, skill curation. Hermes itself is free
Vision	~1,000 analyses; ~2,500 input and ~300 output tokens each	Webcam verification “always”, with bursts of 8–10 frames
GPT-6 Astra	5 consultations, 2 successful; ~40k input and ~30k output each	reasoning max; 2 completed and 1 timed-out consultations documented
P0 phase	~13% of the Hermes volume	Ratio of lines written (10.5k vs 81k)
Total loop + subagents	~1.05 billion input tokens, ~19 million output	Scenarios B and C

Failed attempts (scenario A). Estimated at about 60% of activity: the falsified QMUX hypothesis, v101/v102 cancelled, v121/v122 never flashed, the boot hunt v104–v112, the standby v168–v172. For phase P0 the failure share is estimated at 80%, because the upstream kernel/UEFI/hypervisor route was abandoned entirely.

6.2 Prices applied (dollars per million tokens)

Model	Cheapest condition	Most expensive condition	Notes
DeepSeek V4 Pro	Off-peak: 0.66 input / 0.022 cache / 1.98 output	Peak: 1.32 / 0.044 / 3.96	Peak hours: 01–04 and 06–10 UTC on weekdays; off-peak –50%
DeepSeek V4 Flash (vision, hypothesis)	0.22 / 0.007 / 0.66	0.44 / 0.014 / 1.32	The actual vision model is unknown
GPT-6 Astra	Flex/Batch: 5 / 0.50 / 25	Fast (priority): 20 / 2 / 100	Standard 10 / 1 / 50; above 272K input tokens: 20 / 75
gpt-5.6-luna	Negligible	Negligible	8 test messages
claude-haiku-4-5, OpenRouter	0	0	Failed calls (404 and 402)

In scenario C, besides peak prices and Fast mode, the cache hit rate was reduced by 10 points, because less effective caching is part of the worst conditions.

6.3 Results

Item	A. Minimum	B. Medium	C. Maximum
Orchestrator (main loop)	\$42.77	\$106.92	\$328.68
Hermes overhead	\$3.42	\$8.55	\$26.29
Subagents (fan-outs and delegations)	\$14.58	\$36.45	\$92.45
Webcam vision	\$0.30	\$0.75	\$1.50
GPT-6 Astra	\$1.90	\$4.75	\$19.00
gpt-5.6-luna (test)	≈ \$0	≈ \$0	≈ \$0
P0 phase (unknown agent)	\$3.73	\$18.64	\$54.75
Total	≈ \$67	≈ \$176	≈ \$523




Figure 7 — Estimated costs by scenario and item.




Figure 8 — Percentage breakdown: the main loop accounts for about 60–65%, subagents about 18–22%, orchestration about 5%.




Figure 9 — Cost weight by model or component in scenario B. Astra, despite a unit price 15–25 times higher, accounts for 2.7% because it was used very little.

6.4 Sensitivity and reading of the range




Figure 10 — Sensitivity of the estimate (log scale).

- **The most realistic estimate is close to scenario B.** DeepSeek peak windows cover about 7 hours out of 24 on weekdays and weekends are entirely off-peak, but some documented sessions fall within peak hours (07:15–08:25 and 08:28–12:23 Italian time). Weighting gives about \$190–200.
- **Volume matters more than price.** With halved volumes scenario A drops to about \$33; with volumes 50% higher scenario C rises to about \$784. The defensible range is therefore from about \$35 to about \$800.
- **The decisive variable is the main model.** If the main loop had run on GPT-6 Astra at the standard rate, with the same volumes, that item alone would cost about \$2,460.
- **Unquantified costs:** the markup of the “oneprovider” provider, the web search and scraping APIs used by the fan-outs, the P0 phase tool (which may be a fixed-cost subscription), hardware and human time.

Appendix A — Limitations and next steps

- The session logs (state.db, request_dump_*.json) are cited but not published: with those files, all volume assumptions could be replaced with actual counts per session and per model, and the metrics in section 5 could also become counts.
- Dating by phase uses path, content and folder, because git has no history: assignment errors at the P0/P1 boundary are possible but marginal.
- File classification (third-party, output, code, prose) is automatic, using path rules; counts may vary by a few percentage points.
- The GPT-6 Astra and gpt-5.6-luna models are real and publicly documented; the identity of the “oneprovider” provider and its price list cannot be verified.

Appendix B — External sources

Source	Use
github.com/Saddytech/openwrt-nx679j and release v1.0	Repository identity, description and stated checks
github.com/CAPP-TESTS/openwrt-nx679j	Copy analysed (same commit f9a92fc)
deepseek.ai/pricing · benchlm.ai/deepseek/api-pricing · layer3labs.io/guides/deepseek-pricing · devtk.ai/en/models/deepseek-v4	DeepSeek V4 Pro and Flash prices, peak windows, model capabilities
eesel.ai/blog/gpt-6-astra-pricing · cloudzero.com/blog/gpt-6-pricing · docs.litellm.ai/blog/gpt_6_astra	GPT-6 Astra prices (Standard, Flex/Batch, Fast, long context)
requesty.ai · playground.roboflow.com · cloudprice.net	Existence and characteristics of gpt-5.6-luna and gpt-6-astra
github.com/zhaoxuya520/reverse-skill	Nature of the reverse-skill pack (client-agnostic)
hermes-agent.nousresearch.com	Hermes Agent: MIT licence and custom providers

Appendix C — Key paths in the repository

Path	Relevance
DOCUMENTAZIONE.md (sections “Fase 0d”, “Da chiarire”)	Session provenance, unattributed pre-Hermes phase
08-skill/nx679j-openwrt/references/regole-utente.md	Verbatim user quotes
01-documenti/NX679J-NEW-SESSION-PROMPT.md, MODEM-HANDOFF-20260919.md	Use of GPT-6 Astra via oneprovider, OpenRouter 402 error
01-documenti/wifi-luci/STATO-ATTUALE.md	“DESIGN approvato da Astra”, user authorisations of 23/09
09-albero-originale/experiments-extra/refs/session4-20260920/subagents/	Hermes subagent transcripts
05-bootchain-re/bootloader-re/	P0 phase documents (style, self-corrections)
09-albero-originale/hypervisor-bypass/hyp_attack.c	SM8350 addresses on an SM8450 device
02-sorgenti/wifi-luci/build-v54...v90.py	Last available builders, with absolute paths
DEDUPICA.md	1,450 files removed and references to an unpublished revision archive